

# Generative-AI Powered Inference with gpi-pack

Kentaro Nakamura


Harvard University

Oxford Computational Political Science Group

May 20, 2025

Joint work with Kosuke Imai

gpi-pack: <https://gpi-pack.github.io/>



Search docs

GETTING STARTED

- Installation
- What's GPI?
- How to use GPU

DATA GENERATION

- Generating Texts with LLaMa3
- Generating Texts with Other LLMs

BASIC OPERATION

- Text-As-Treatment

ADVANCED OPERATIONS

- Hyperparameter Tuning
- Customizing Your Analysis
- When LLM is too big

 / GPI: Generative-AI Powered Inference

## GPI: Generative-AI Powered Inference

pypi [v0.1.0](#) python [3.8](#) | [3.9](#) | [3.10](#) | [3.11](#)



**gpi\_pack** is a Python library for the statistical inference powered by Generative Artificial Intelligence (AI). It provides a set of tools and utilities for performing statistical inference using the internal representation of the Generative AI models. The library is designed to be easy to use and flexible, allowing users to perform a wide range of statistical analyses.

### Note

We released **gpi\_pack** version 0.1.0 on February 27th, 2025. This is the first version of the package, and we currently only support the setting of Text-as-Treatment based on [our paper](#). We have been working hard to make this package as useful and user-friendly as possible. If you have any feedback or suggestions, please feel free to reach out to [the maintainer](#).

# Computing Environment: How to use GPU

- `gpi-pack` is built upon PyTorch that supports GPU acceleration

# Computing Environment: How to use GPU

- `gpi-pack` is built upon PyTorch that supports GPU acceleration
- GPU is not essential, but it significantly speeds up computation

# Computing Environment: How to use GPU

- `gpi-pack` is built upon PyTorch that supports GPU acceleration
- GPU is not essential, but it significantly speeds up computation
  
- Two ways to use GPU

# Computing Environment: How to use GPU

- `gpi-pack` is built upon PyTorch that supports GPU acceleration
- GPU is not essential, but it significantly speeds up computation
- Two ways to use GPU
  - 1 Local GPU Computing

# Computing Environment: How to use GPU

- `gpi-pack` is built upon PyTorch that supports GPU acceleration
- GPU is not essential, but it significantly speeds up computation
- Two ways to use GPU
  - 1 Local GPU Computing
  - 2 Cloud GPU computing platforms (e.g., AWS, Google Colab)

# Computing Environment: How to use GPU

- gpi-pack is built upon PyTorch that supports GPU acceleration
- GPU is not essential, but it significantly speeds up computation
- Two ways to use GPU
  - 1 Local GPU Computing
  - 2 Cloud GPU computing platforms (e.g., AWS, Google Colab)
- Check if GPU is available:

```
1 import torch
2 print(torch.cuda.is_available())
```



Google Colab: <https://colab.research.google.com/>

- Free, cloud Jupyter notebook environment offering access to GPUs

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**
  - ① Create a notebook / Upload notebook to Google Drive

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**
  - 1 Create a notebook / Upload notebook to Google Drive
  - 2 Go to Runtime and Change Runtime Type

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**
  - 1 Create a notebook / Upload notebook to Google Drive
  - 2 Go to Runtime and Change Runtime Type
  - 3 Select GPU as the hardware accelerator

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**
  - 1 Create a notebook / Upload notebook to Google Drive
  - 2 Go to Runtime and Change Runtime Type
  - 3 Select GPU as the hardware accelerator
- Free tier has limits (e.g., 12 hours per session)

- Free, cloud Jupyter notebook environment offering access to GPUs
- **How to use:**
  - 1 Create a notebook / Upload notebook to Google Drive
  - 2 Go to Runtime and Change Runtime Type
  - 3 Select GPU as the hardware accelerator
- Free tier has limits (e.g., 12 hours per session)
- Colab Pro/Pro+ offers more compute and access to premium GPUs

# Three Steps of Implementation

- 1 Generating internal representation with Generative-AI



# Three Steps of Implementation

- ① Generating internal representation with Generative-AI
  - Text: Open-Source LLM (e.g., LLaMa, Gemma)

# Three Steps of Implementation

- ① Generating internal representation with Generative-AI
  - Text: Open-Source LLM (e.g., LLaMa, Gemma)
  - Image: Diffusion model (e.g., Stable Diffusion)

# Three Steps of Implementation

- ① Generating internal representation with Generative-AI
  - Text: Open-Source LLM (e.g., LLaMa, Gemma)
  - Image: Diffusion model (e.g., Stable Diffusion)
- ② Hyperparameter tuning of nuisance functions for the outcome model

# Three Steps of Implementation

- 1 Generating internal representation with Generative-AI
  - Text: Open-Source LLM (e.g., LLaMa, Gemma)
  - Image: Diffusion model (e.g., Stable Diffusion)
- 2 Hyperparameter tuning of nuisance functions for the outcome model
- 3 Estimate causal effects

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models



# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)
    - Model quantization often helps

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)
    - Model quantization often helps
  - ② Using API that allows users to access internal states

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)
    - Model quantization often helps
  - ② Using API that allows users to access internal states
    - National Deep Inference Fabric (NDIF) offers the package `nnsight`

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)
    - Model quantization often helps
  - ② Using API that allows users to access internal states
    - National Deep Inference Fabric (NDIF) offers the package `nnsight`
    - Website: <https://nnsight.net/>

# STEP1: Generating Internal Representation with GenAI

- Two prompting strategies
  - ① **Create**: create texts / images from scratch
  - ② **Reuse**: repeating the same texts and images as the inputs
- Two options to use deep generative models
  - ① Deploying models locally
    - Requires GPU, computationally heavy
    - Impossible to deploy the heavy model (e.g., LLaMa-495B)
    - Model quantization often helps
  - ② Using API that allows users to access internal states
    - National Deep Inference Fabric (NDIF) offers the package `nnsight`
    - Website: <https://nnsight.net/>
    - Plan to release the functionality for `nnsight` soon



## Example: Deploying models locally (1)

- Find your favorite model from <https://huggingface.co/>

```
1 from transformers import AutoTokenizer, AutoModelForCausalLM
2 import torch
3
4 checkpoint = 'meta-llama/Meta-Llama-3.1-8B-Instruct'
5
6 tokenizer = AutoTokenizer.from_pretrained(checkpoint)
7 model = AutoModelForCausalLM.from_pretrained(
8     checkpoint,
9     device_map="auto",
10    torch_dtype=torch.float16
11 )
```

## Example: Deploying models locally (1)

- Find your favorite model from <https://huggingface.co/>
  - Typically, the documentation shows how to deploy the model

```
1 from transformers import AutoTokenizer, AutoModelForCausalLM
2 import torch
3
4 checkpoint = 'meta-llama/Meta-Llama-3.1-8B-Instruct'
5
6 tokenizer = AutoTokenizer.from_pretrained(checkpoint)
7 model = AutoModelForCausalLM.from_pretrained(
8     checkpoint,
9     device_map="auto",
10     torch_dtype=torch.float16
11 )
```

## Example: Deploying models locally (2)

- Use `extract_and_save_hidden_states` function in `gpi_pack`

```
1 from gpi_pack.llm import extract_and_save_hidden_states
2
3 prompts = [
4     'Create a biography of a politician named John Doe',
5     'Create a biography of a politician named Jane Smith',
6 ]
7
8 extract_and_save_hidden_states(
9     prompts = prompts,
10    output_hidden_dir = <YOUR HIDDEN DIR>,
11    save_name = <YOUR SAVE NAME>,
12    tokenizer = tokenizer,
13    model = model,
14    task_type = "create"
15 )
```

## Example: Deploying models locally (3)

- To repeat the input texts, set `task_type == "repeat"`

```
1 from gpi_pack.llm import extract_and_save_hidden_states
2
3 prompts = [
4     'The Airports Commission, an independent body
5     established...',
6     'History show us that most large infrastructure projects
7     ...',
8 ]
9
10 extract_and_save_hidden_states(
11     prompts = prompts,
12     output_hidden_dir = <YOUR HIDDEN DIR>,
13     save_name = <YOUR SAVE NAME>,
14     tokenizer = tokenizer,
15     model = model,
16     task_type = "repeat"
```

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate



## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size
  - dropout rate

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size
  - dropout rate
  - size and width

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size
  - dropout rate
  - size and width
- Optuna offers efficient hyperparameter optimization

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size
  - dropout rate
  - size and width
- Optuna offers efficient hyperparameter optimization
  - `gpi_pack` offers the class `TarNetHyperparameterTuner`

## STEP2: Hyperparameter Tuning

- The performance of the estimator is sensitive to hyperparameters
  - Fitting of deconfounder is especially important
- Hyperparameters
  - learning rate
  - batch size
  - dropout rate
  - size and width
- Optuna offers efficient hyperparameter optimization
  - `gpi_pack` offers the class `TarNetHyperparameterTuner`
  - Only need to specify data and the choice of hyperparameters

## Example: Hyperparameter Tuning with Optuna

```
1 from gpi_pack.TarNet import TarNetHyperparameterTuner
2 import optuna
3
4 obj = TarNetHyperparameterTuner(
5     # Data
6     T = df['TreatmentVar'].values,
7     Y = df['OutcomeVar'].values,
8     R = hidden_states,
9     # Hyperparameters
10    learning_rate = [1e-4, 1e-5],
11    dropout = [0.1, 0.2],
12    architecture_y = ["[200, 1]", "[100,1]"],
13    architecture_z = ["[1024]", "[2048]"]
14 )
15
16 # Hyperparameter tuning with Optuna
17 study = optuna.create_study(direction='minimize')
18 study.optimize(obj.objective, n_trials=100)
```

## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect



## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect
- `estimate_ate_k` function deals with all the procedures

## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect
- `estimate_ate_k` function deals with all the procedures
  - Propensity score function needs to be Lipschitz continuous

## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect
- `estimate_ate_k` function deals with all the procedures
  - Propensity score function needs to be Lipschitz continuous
  - Default choice is neural network with spectral normalization

## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect
- `estimate_ate_k` function deals with all the procedures
  - Propensity score function needs to be Lipschitz continuous
  - Default choice is neural network with spectral normalization
- You should always plot the estimated propensity score distributions


## STEP3: Estimate Treatment Effect

- K-fold cross-fitting to estimate the treatment effect
- `estimate_ate_k` function deals with all the procedures
  - Propensity score function needs to be Lipschitz continuous
  - Default choice is neural network with spectral normalization
- You should always plot the estimated propensity score distributions
  - For text/image-as-treatment, it works as a diagnosis of separability assumption

## Example: Estimate Treatment Effect

```
1 # estimate treatment effects
2 ate, se = estimate_k_ate(
3     R= hidden_states,
4     Y= df['OutcomeVar'].values,
5     T= df['TreatmentVar'].values,
6     K=2, #K-fold cross-fitting
7     lr = 2e-5, #learning rate
8     architecture_y = [200, 1], #outcome model architecture
9     architecture_z = [2048], #deconfounder architecture
10    plot_propensity = True, #visualize propensity scores
11 )
```

## Example: Propensity Score Distributions



# Conclusion

- `gpi-pack` is an open-source software to implement GPI



# Conclusion

- gpi-pack is an open-source software to implement GPI
- Website: <https://gpi-pack.github.io/>

# Conclusion

- gpi-pack is an open-source software to implement GPI
- Website: <https://gpi-pack.github.io/>
- If you have any question, find bugs, or have any suggestions,

# Conclusion

- gpi-pack is an open-source software to implement GPI
- Website: <https://gpi-pack.github.io/>
- If you have any question, find bugs, or have any suggestions,
  - 1 Please open an issue on GitHub  
([https://github.com/gpi-pack/gpi\\_pack](https://github.com/gpi-pack/gpi_pack))

# Conclusion

- gpi-pack is an open-source software to implement GPI
- Website: <https://gpi-pack.github.io/>
- If you have any question, find bugs, or have any suggestions,
  - 1 Please open an issue on GitHub  
([https://github.com/gpi-pack/gpi\\_pack](https://github.com/gpi-pack/gpi_pack))
  - 2 Email me at knakamura [at] g.harvard.edu